

Google Cloud Professional Cloud Network Engineer Exam Prep Notes by

Ammett

V2



Google Cloud Professional Cloud Network Engineer

Exam Prep Sheet by Ammett

This is my updated 2021 guide based on my preparation for the exam. References from Google Docs and other sources.

V2: 03-2021







- 1- [Best practices for enterprise organizations](#)
- 2- [VPC Overview](#)
- 3- [Alias IP](#)
- 4- [VPC Network Peering](#)
- 5- [Shared VPC](#)

- 6- [Choosing a load balancer](#)
- 7- [Cloud CDN Overview](#)
- 8- [Choosing a VPN option](#)
- 9- [Cloud Router](#)
- 10- [Direct Peering](#)
- 11- [Carrier peering](#)

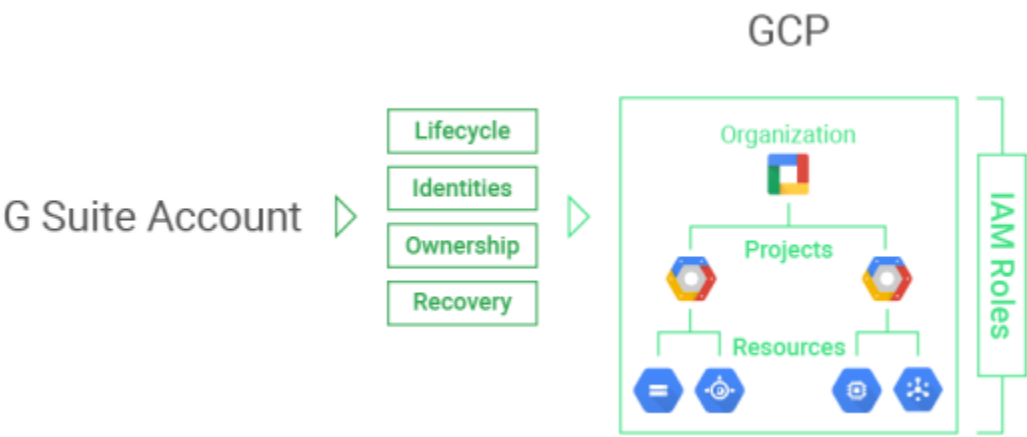
- 12- [Dedicated Interconnect](#)
- 13- [Partner Interconnect](#)
- 14- [Creating a VPC-native Cluster](#)
- 15- [Private Cluster Kubernetes](#)
- 16- [Firewall Rules Logging](#)

- 17- [Cloud DNS](#)
- 18- [Networking Kubernetes](#)
- 19- [URL_Map](#)
- 20- [Load balancer health checks](#)

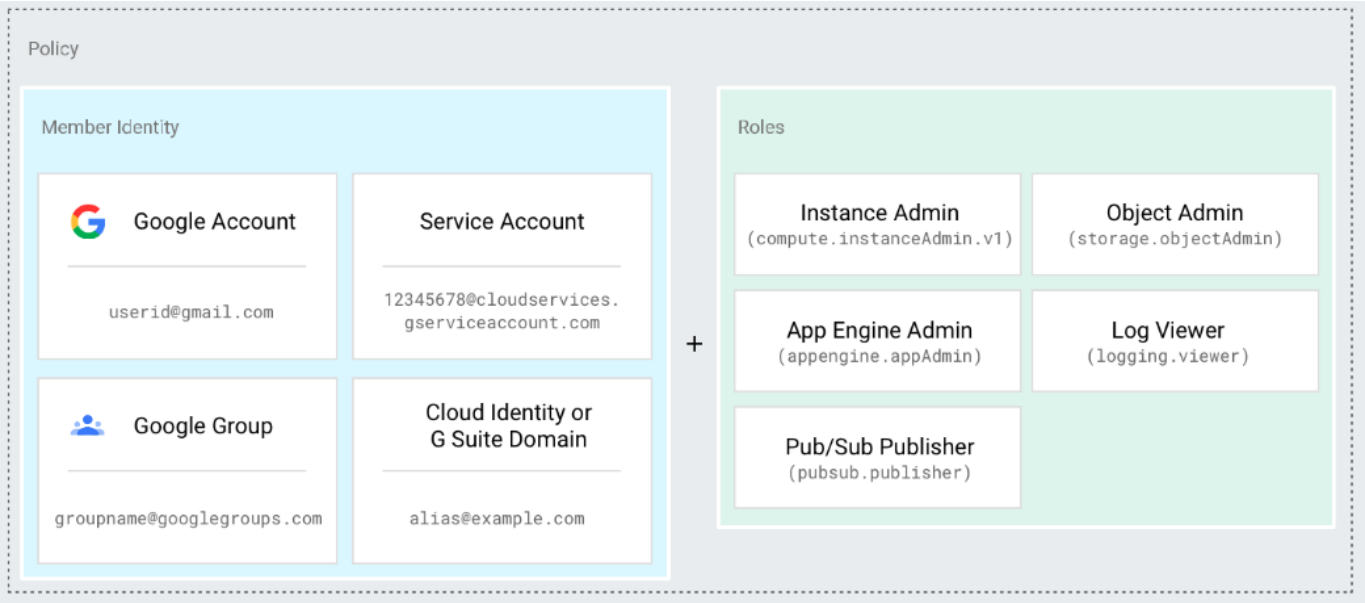


Organisation Structures 	What it is Resources are organized hierarchically. This allows you to map your enterprise's operational structure to GCP, and to manage access control and permissions for groups of related resources.	Key points 1- Flow (Organisation, Folders, Projects, Resources) 2- Where to manage permissions for groups, department, entire organisation, etc	What you should know 1- Permissions level necessary to do certain functions 2- Domains, Groups, G Suite domain, Super users.	Review documents Cloud Platform hierarchy	Video Hierarchy	My experience This area is fundamental please understand how to control to get the separation and security in your domain.
Cloud IAM 	What it is IAM which lets you manage access control by defining who (identity) has what access (role) for which resource.	Key points 1- Best way to manage (use groups) 2- Roles (primitive, predefined & custom) 3- Roles necessary to do certain functions (network, security, IAM, cloud storage)	What you should know 1- Permissions level necessary 2- Permission errors 3- How & when to create custom roles 4- Service account permissions	Review documents Cloud IAM overview	Video Cloud IAM Best practices for identity	My experience IAM on a networking exam? Yes, know it well because it will come. Knowing the roles necessary for certain actions may help if you can figure it out.
CIDR RFC-1918 	What it is You can choose any private RFC 1918 CIDR block for the primary IP address range of the subnet	Key points 1- The 4 Reserved Address (network, gateway, google reserved, broadcast) 2- How to assign your own range	What you should know 1- How to assign static internal IP 2- How to change IP	Review documents IP Addresses Reserve Internal IP	Video Networking with IP Address	My experience Some form of RFC-1918 will come. Keep in mind what is reserved, auto-mode RFC 1918 addresses.
External IP 	What it is These are routable on the public internet and allow you access to the internet.	Key points 1- This is optional 2- Default is ephemeral-these change 3- Static can be assigned 3- How to create static external IP	What you should know 1- Charged if not attached to VM 2- How to change ephemeral IP to another ephemeral IP	Review documents Reserve External IP		My experience These can appear but shouldn't be too difficult to handle
Subnet Types 	What it is Subnets are used to separate resources and control communication between tiers. Access can be controlled via routes and firewalls	Key points 1- Default (automatically generated with a project) they have default firewall rules and a subnet in every region 2- Auto-mode- automatically creates a subnet in every region (the default subnet is an auto mode subnet) IP range 10.128.0.0/9	What you should know 1- Custom is fully user controlled 2- Avoid overlapping ranges 3- You can convert from auto to custom (one way). Things can get affected. 4- You can increase range not decrease		Video Create Custom Subnet	My experience Take note of this area. CIDR block host availability for VPC and also in Kubernetes.
Private Access 	What it is Allows VM with internal (RFC 1918) IP addresses to reach certain APIs and services without internet access.	Key points 1- No public IP address 2- Enabled on subnet 3- Default route	What you should know 1- Services that support Private access 2- Default route 0.0.0.0/0 next hop "default internet gateway" or custom routes 199.36.153.4/30 or 199.36.153.8/30 nexthop "default internet gateway"	Review documents Private Google Access	Video Access GCP and 3rd party services privately	My experience This is a key topic . Especially what services are supported and how to set up.

IAM example 1



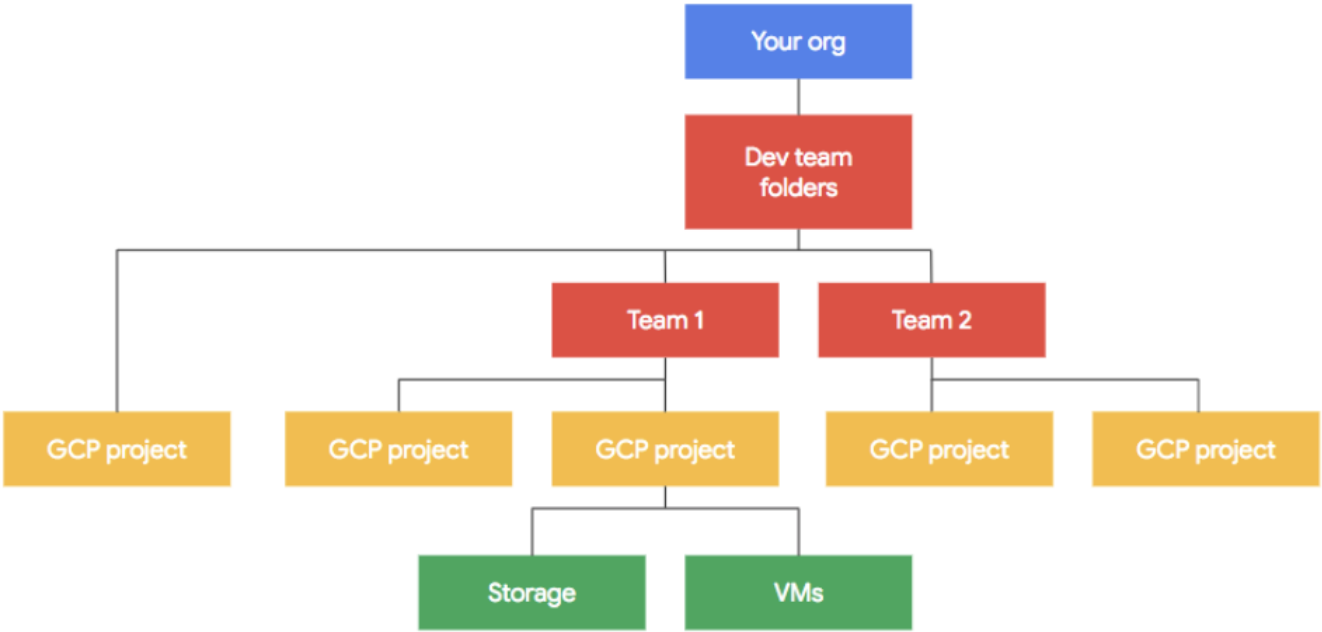
IAM example 2



Reserved range example

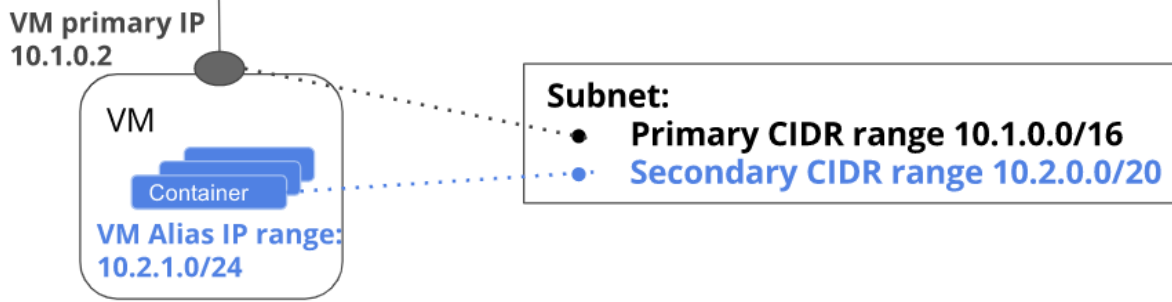
Reserved Address	Description	Example
Network	First address in the primary IP range for the subnet	10.1.2.0 in 10.1.2.0/24
Default gateway	Second address in the primary IP range for the subnet	10.1.2.1 in 10.1.2.0/24
Second-to-last address	Second-to-last address in the primary IP range for the subnet that is reserved by GCP for potential future use	10.1.2.254 in 10.1.2.0/24
Broadcast	Last address in the primary IP range for the subnet	10.1.2.255 in 10.1.2.0/24

Hierarchy Flow



Private Service 	What it is The private connection enables VM in your VPC network and the services that you access to communicate exclusively by using internal (RFC 1918) IP addresses.	Key points 1- External IP addresses are not required or used 2- Service producers network 3- Private IP 4- Cloud SQL supports this	What you should know 1- Works via peering from customer to service producer network 2- Must define CIDR range for services. 3- Connect within same region	Review documents Config private service access	Video Access GCP and 3rd party services privately Labs Connecting to Cloud SQL: Private IP Multiple VPC VPC network	My experience Know difference between Private services and Private access. Know which services use which.
Alias IP 	What it is Alias IP ranges let you assign ranges of internal IP addresses as aliases to a (VM) nic. This is useful if you have multiple services running on a VM and you want to assign each service a different IP address. Alias IP ranges also work with GKE Pods.	Key points 1- Main address from primary CIDR 2- Alias can be from main CIDR or 3- Alias IP can be from secondary ranges.	What you should know 1- Use of alias IP ranges does not require secondary subnet ranges. These secondary subnet ranges merely provide an organizational tool.	Review documents Alias IP Configuring Alias IP		My experience Pay attention here. Alias IP can be used in VM and Kubernetes.
VPC 	What it is A VPC network is your virtual network in the cloud just like an on premise physical network or data centre or office network.	Key points 1- VPC are global SDN 2- How to get traffic flowing 3- Using RFC 1918 subnets 4- Internal and external access	What you should know 1- Internal and external access 2- Controlling access and firewalls 3- How to Connect VPC together (peering or sharing)	Review documents VPC Overview	Video VPC Deep Dive	My experience Core area. Let me put it like this; If you do not understand all of the elements of a VPC; then don't do the exam.
Routes 	What it is These define the paths network traffic takes from a VM instance to other destinations. These destinations can be inside or outside of your VPC.	Key points 1- The route table is defined at network level 2- The routing to next hop where should the next hop be	What you should know 1- Type (system and custom) 2- Default route & Subnet route 3- Static and Dynamic routes 4- Routing order	Review documents Routes in GCP	Video Cloud Router	My experience You cannot have networking without routes. (Static, dynamic, subnet, custom, default, import, export)
Cloud Router 	What it is This enables you too dynamically exchange routes between (VPC) and on-premises networks by using Border Gateway Protocol (BGP).	Key points 1- Cloud Router automatically learns new subnets in your VPC network and announces them to your on-premises network	What you should know 1- Global dynamic routing 2- Regional dynamic routing	Review documents Cloud Router		My experience Another critical area. Know how these are setup. Has lot of small parts get familiar.
BGP 	What it is Border Gateway Protocol is a protocol that manages how packets are routed across the internet through the exchange of routing and reachability information between edge routers .	Key points 1- The ASN number range (64512 - 65534, 4200000000 – 4294967294) 2- IP range used 169.254.0.0/16	What you should know 1- MED (route priority) 2- What can be configured without BGP 3- Troubleshooting	Review documents Establishing BGP sessions Troubleshooting Cloud Router		My experience A question or 3 may come on BGP. Know what is required, problems and how it works.
Firewall 	What it is Allow or deny traffic to and from your virtual machine (VM) etc, based on configurations you specify.	Key points 1- How they work (Stateful) & Scope 2- Implied rules 3- Default rules	What you should know 1- How to restrict traffic 2- Use of tag 2- Use of service account	Review documents Firewalls	Video Firewalls Network and security telemetry	My experience You can't allow everything on your network so expect a few firewall questions in the networking exam also.
Firewall logging 	What it is Firewall Rules Logging allows you audit, verify, and analyze the effects of your firewall rules.	Key points 1- Individually enabled 2- Supported for TCP & UDP only 3- Cannot enable on implied or default rules	What you should know 1- Troubleshooting viewing (Log entries missing, cannot view logs, where to apply logs)	Review documents Firewall Logging		My experience You should have an idea where to look, what rules are logged, priorities and how to fix.

Alias IP image example



Trouble shooting logs

Log entries missing

Possible cause: Connections might not match the firewall rule you expect

Verify that the firewall rule you expect is in the list of applicable firewall rules for an instance. Use the GCP Console to view details for the relevant instance, then click the **View details** button in the *Network interfaces* section on its *VM instance details* page. Inspect applicable firewall rules in the *Firewall rules and routes details* section of its *Network interface details* page.

Review the [firewall rules overview](#) to make sure you have created your firewall rules correctly.

You can use [tcpdump](#) on the VM to determine if connections it sends or receives have addresses, ports, and protocols that would match the firewall you expect.

Possible cause: A higher priority rule with firewall rules logging disabled might apply

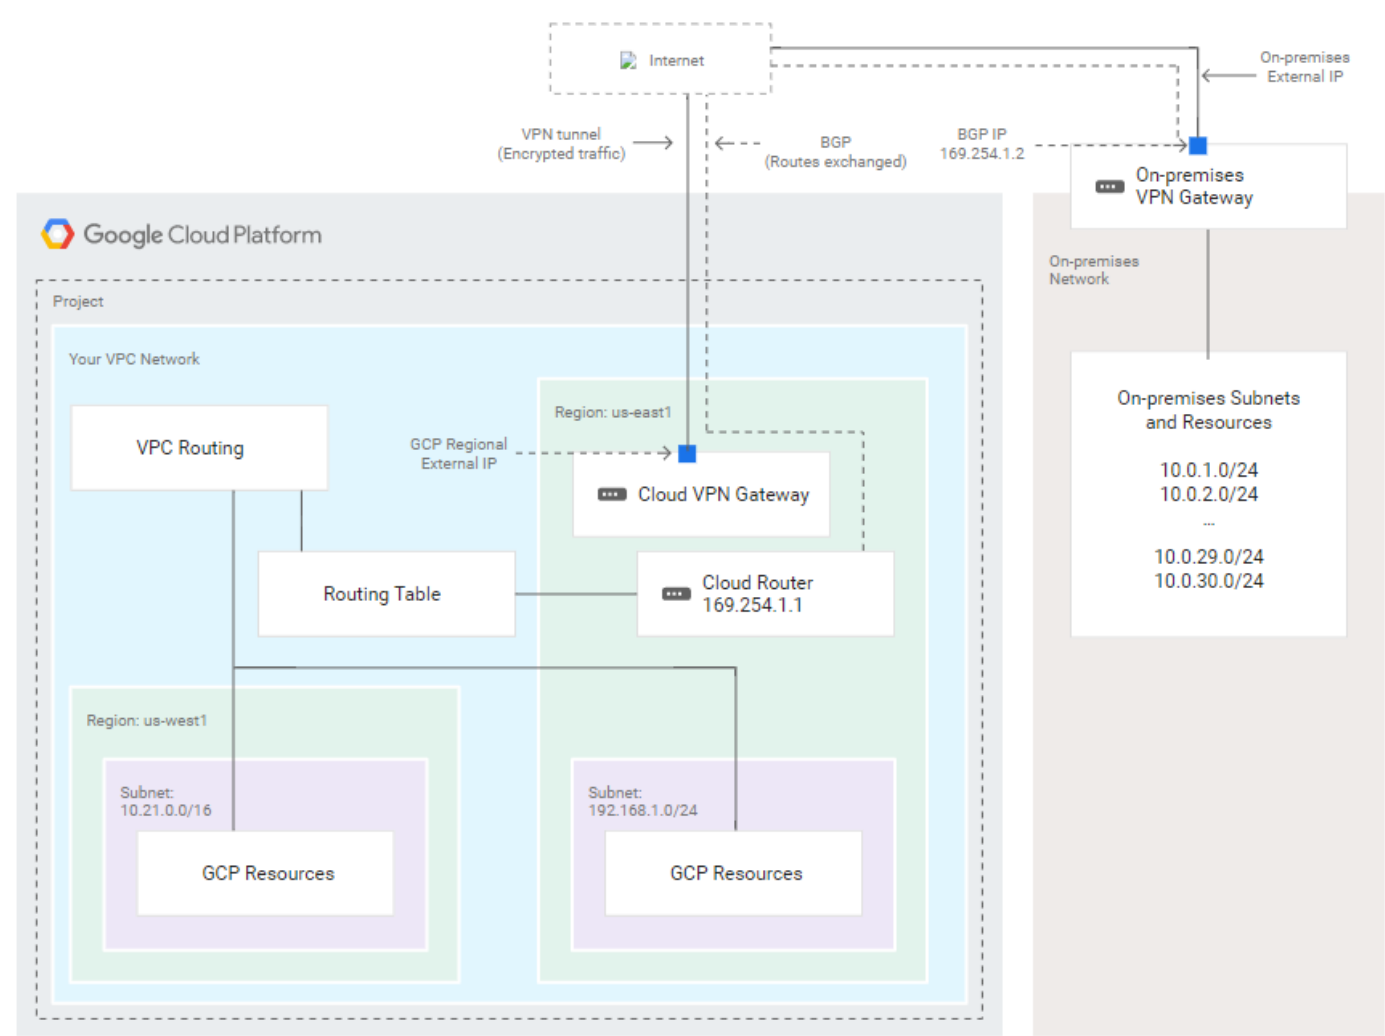
Firewall rules are [evaluated according to their priorities](#). From the perspective of a VM instance, only one firewall rule applies to the traffic.













A rule that you think would be the highest priority applicable rule might not actually be the highest priority applicable rule. A higher priority rule that does not have logging enabled might apply instead.

To troubleshoot, you can temporarily enable logging for *all* possible firewall rules applicable to a VM. Use the GCP Console to view details for the relevant VM, then click the **View details** button in the *Network interfaces* section on its *VM instance details* page. Inspect applicable firewall rules in the *Firewall rules and routes details* section of its *Network interface details* page, and identify your custom rules in that list. Temporarily enable logging for all of those custom firewall rules.

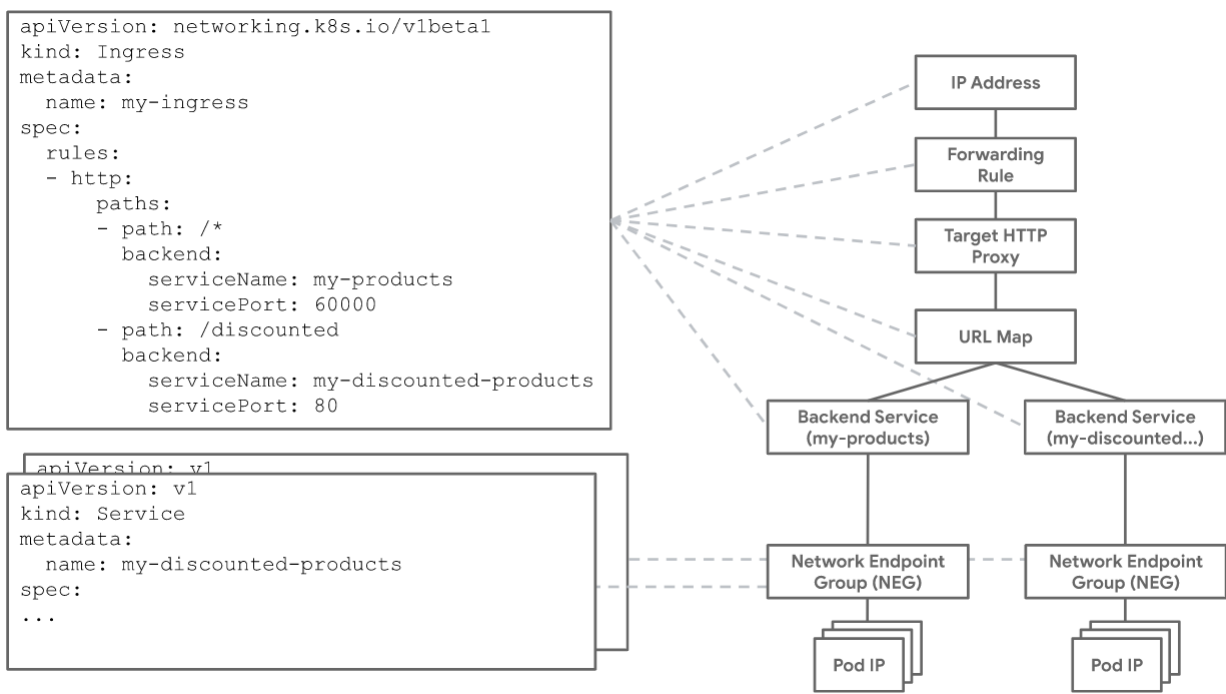
With logging enabled, you can identify the applicable rule. Once identified, be sure to disable logging for all rules that do not actually need it.

Cloud Router for VPNs with VPC network

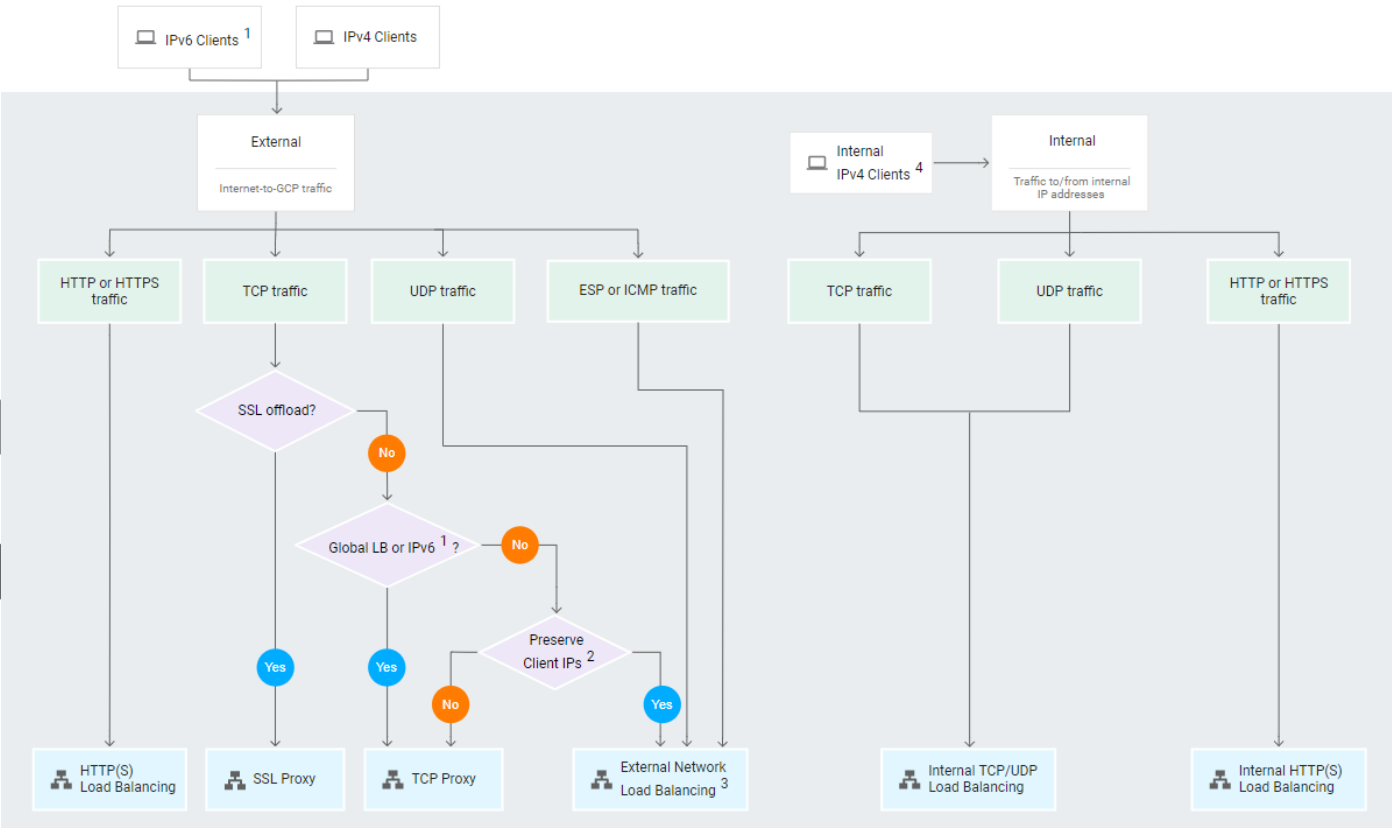


<div>HTTP(S) Load balancer</div> <div></div>	<div>SSL Proxy</div> <div></div>	<div>TCP Proxy</div> <div></div>	<div>Network Load balancer</div> <div></div>	<div>Internal load balancer</div> <div></div>	<div>Kubernetes Load balancing</div> <div></div>	<div>Review documents</div> <div>Choosing a load balanced</div> <div>Troubleshooting health</div> <div>HTTPS logging</div> <div>Kubernetes HTTP(s) LB ingress</div>
<div>What it is</div> <div>Load balancer for HTTP(S) traffic, global, external, 80 or 8080 on 443.</div>	<div>What it is</div> <div>Load balancer for TCP with SSL offload, global, external. (25, 43, 110, 143,195, 443, 465, 587, 700, 993, 995, 1883, and 5222)</div>	<div>What it is</div> <div>Load balancer for TCP without SSL, global, external. (25, 43, 110, 143,195, 443, 465, 587, 700, 993, 995, 1883, and 5222)</div>	<div>What it is</div> <div>Load balancer for TCP/UDP no SSL offload, regional, external. (any port)</div>	<div>What it is</div> <div>Load balancer for TCP /UDP regional, Internal traffic (any port)</div>	<div>What it is</div> <div>This allows you balance between you application running in Kubernetes</div>	<div>Setting up HTTP Ingress LB</div> <div>Video</div> <div>Cloud Load balancers</div>
<div>What you should know</div> <div>1- Scope global</div> <div>2- HTTPS traffic</div> <div>3- Health checks</div>	<div>What you should know</div> <div>1- Scope Global</div> <div>2- Non HTTPS traffic with SSL termination</div>	<div>What you should know</div> <div>1- Scope Global</div> <div>2- TCP/UDP traffic</div> <div>3- Health checks</div>	<div>What you should know</div> <div>1- Scope regional</div> <div>2- TCP/UDP traffic</div> <div>3- Health checks</div>	<div>What you should know</div> <div>1- Scope Regional</div> <div>2- Internal TCP/UDP traffic</div>	<div>What you should know</div> <div>1- How it works</div> <div>2- Connections points</div> <div>3- Type of LB supported (HTTPS-Ingress, Internal, External)</div>	<div>My experience</div> <div>Loads and loads</div> <div>of variation on this area. (Global vs Regional, External vs Internal, Traffic type, VoIP, TFTP, IP, TCP, UDP). Understand health checks checks. For Kubernetes understand connection points of load balancers. If you don't understand these don't do the exam.</div>
<div>Key Points</div> <div>1- Services that need HTTPS Load balancing</div>	<div>Key Points</div> <div>1- SSL termination</div>	<div>Key Points</div> <div>1- Scope global</div>	<div>Key Points</div> <div>1- Scope regional</div>	<div>Key Points</div> <div>1- Scope regional</div>	<div>Key Points</div> <div>1- What IP you connect to</div> <div>2- HTTPS traffic</div>	
<div>DDoS</div> <div></div>	<div>URL-Mapping</div> <div></div>	<div>Managed Instance Groups</div> <div></div>	<div>Unmanaged Instance Groups</div> <div></div>	<div>Canary Deployments</div> <div></div>	<div>Rolling Deployments</div> <div></div>	<div>Review documents</div> <div>Rolling Updates</div> <div>Managed instances</div> <div>Unmanaged instances</div> <div>URL Map</div>
<div>What it is</div> <div>A (DDoS) attack is a malicious attempt to disrupt normal traffic to a targeted service or network by overwhelming the target infrastructure with a flood of Internet traffic.</div>	<div>What it is</div> <div>Google Cloud Platform HTTP(S) load balancers use a URL map to direct incoming requests to backend services and backend buckets.</div>	<div>What it is</div> <div>A managed instance group contains identical instances that you can manage as a single entity in a single zone.</div>	<div>What it is</div> <div>Unmanaged instance groups are collections of instances that are not necessarily identical and do not share a common instance template.</div>	<div>What it is</div> <div>A canary update is an update that is applied to a partial number of instances in the instance group.</div>	<div>What it is</div> <div>A rolling update is an update that is gradually applied to all instances in an instance group until all instances have been updated</div>	<div>Video</div> <div>Highly available deployments</div> <div>Labs</div> <div>Create Internal LB</div>
<div>What you should know</div> <div>1- How to prevent with GCP tools</div>	<div>What you should know</div> <div>1- How to configure</div> <div>2- It works with HTTPS LB's</div>	<div>What you should know</div> <div>1- Global</div> <div>2- TCP/UDP traffic</div> <div>3- Health checks</div>	<div>What you should know</div> <div>1- When to use.</div> <div>2- Different template.</div>	<div>What you should know</div> <div>1- Applies to a defined amount or % of host</div>	<div>What you should know</div> <div>1- Applies to 100% of the target as defined</div> <div>2- You can configure time etc</div>	
<div>Key Points</div> <div>1- Traffic controlling tools is necessary</div>	<div>Key Points</div> <div>1- Hostname and path</div> <div>2- Characters / an *</div>	<div>Key Points</div> <div>1- Managed instance groups support Autoscaling, load balancing, rolling updates, autohealing, and more.</div>	<div>Key Points</div> <div>1- Unmanaged groups do not create, delete, or scale the number of instances in the group.</div>	<div>Key Points</div> <div>1- Understand when to use for minimization of application performance issues</div>	<div>Key Points</div> <div>1- Understand when to use and impact on application performance</div>	<div>My experience</div> <div>All these area combined made for some VERY challenging questions. Kubernetes is well represented, learn networking, subnetting and load balancing well.</div>


Kubernetes networking example









Choosing a Load balancer diagram



URL Map example

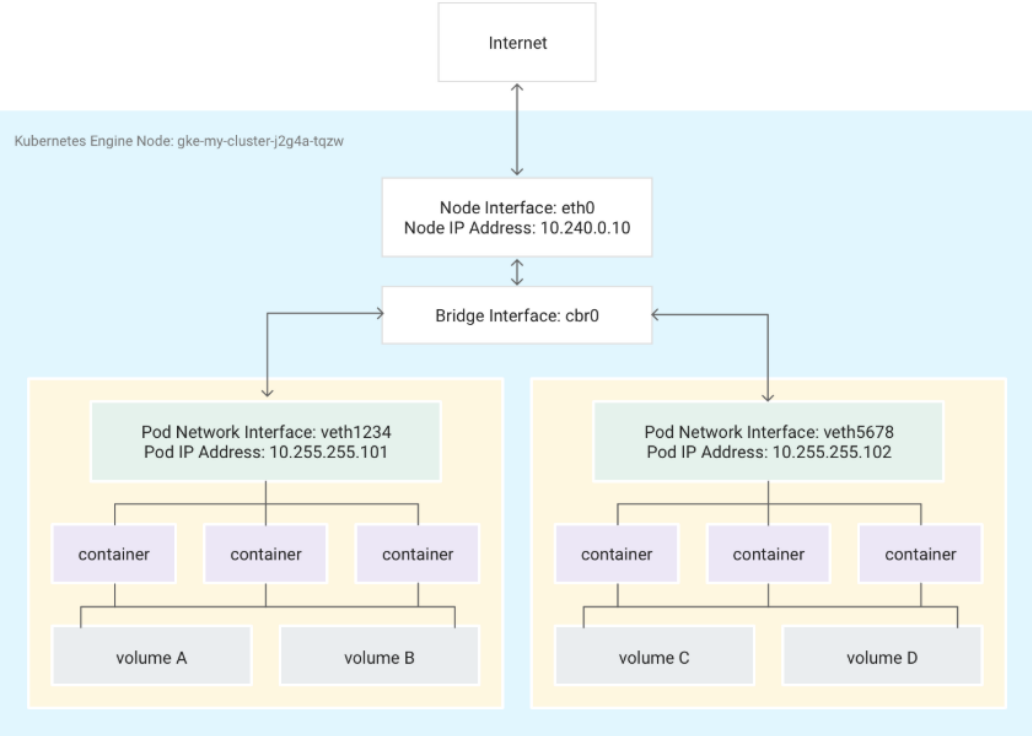














Google Kubernetes Engine	Cluster	Node	Pods	IP tables	Kubernetes subnetting	Review documents
						Networking in Kubernetes
What it is GKE provides a managed environment for deploying, managing, and scaling your containerized applications using Google infrastructure.	What it is A cluster consists of at least one cluster master and multiple nodes. These master and node machines run the Kubernetes cluster orchestration system	What it is They are the worker machines that run your containerized applications and workloads. Each node is managed from the master.	What it is Pods are the smallest, most basic deployable objects in Kubernetes. Pods contain one or more <i>containers</i> ,	What it is Kube-proxy manages the iptables rules on the node.	What it is IP addresses are used for Pods, Nodes and services. The IP subnetting scheme must take into consideration enough for expansion.	Video Deep dive Into Kubernetes Networking
What you should know 1- IP (Cluster, Pod, Node) 2- IP allocation to (nodes, pods, services) 3- Health checks	What you should know 1- Kubernetes object all run on top a cluster 2- Cluster master runs control plane, API server, scheduler and resource controllers.	What you should know 1- You can run a maximum of 110 Pods on a node with a /24 range. 2- Node run kubelet and services to support Docker containers. 3- IP assigned from primary range.	What you should know 1- Pods are ephemeral. 2- Pods do not "heal" or repair themselves. 3- Containers in a pod communicate via local host 4- IP assigned to Virtual NIC in the pod's network namespace.	What you should know 1- Facilitate forwarding within a cluster. 2- These differ from one scenario to the other	What you should know 1- How to assigned based on network requirement (Node, Pod, Services/Cluster IP) 2- Know subnet host count and restrictions 3- Node get IP from primary range, Pod and services from secondary range.	My experience Kubernetes is represented as it should be on this exam. Pay attention to the networking components, subnetting and structure.

Subnetting guide

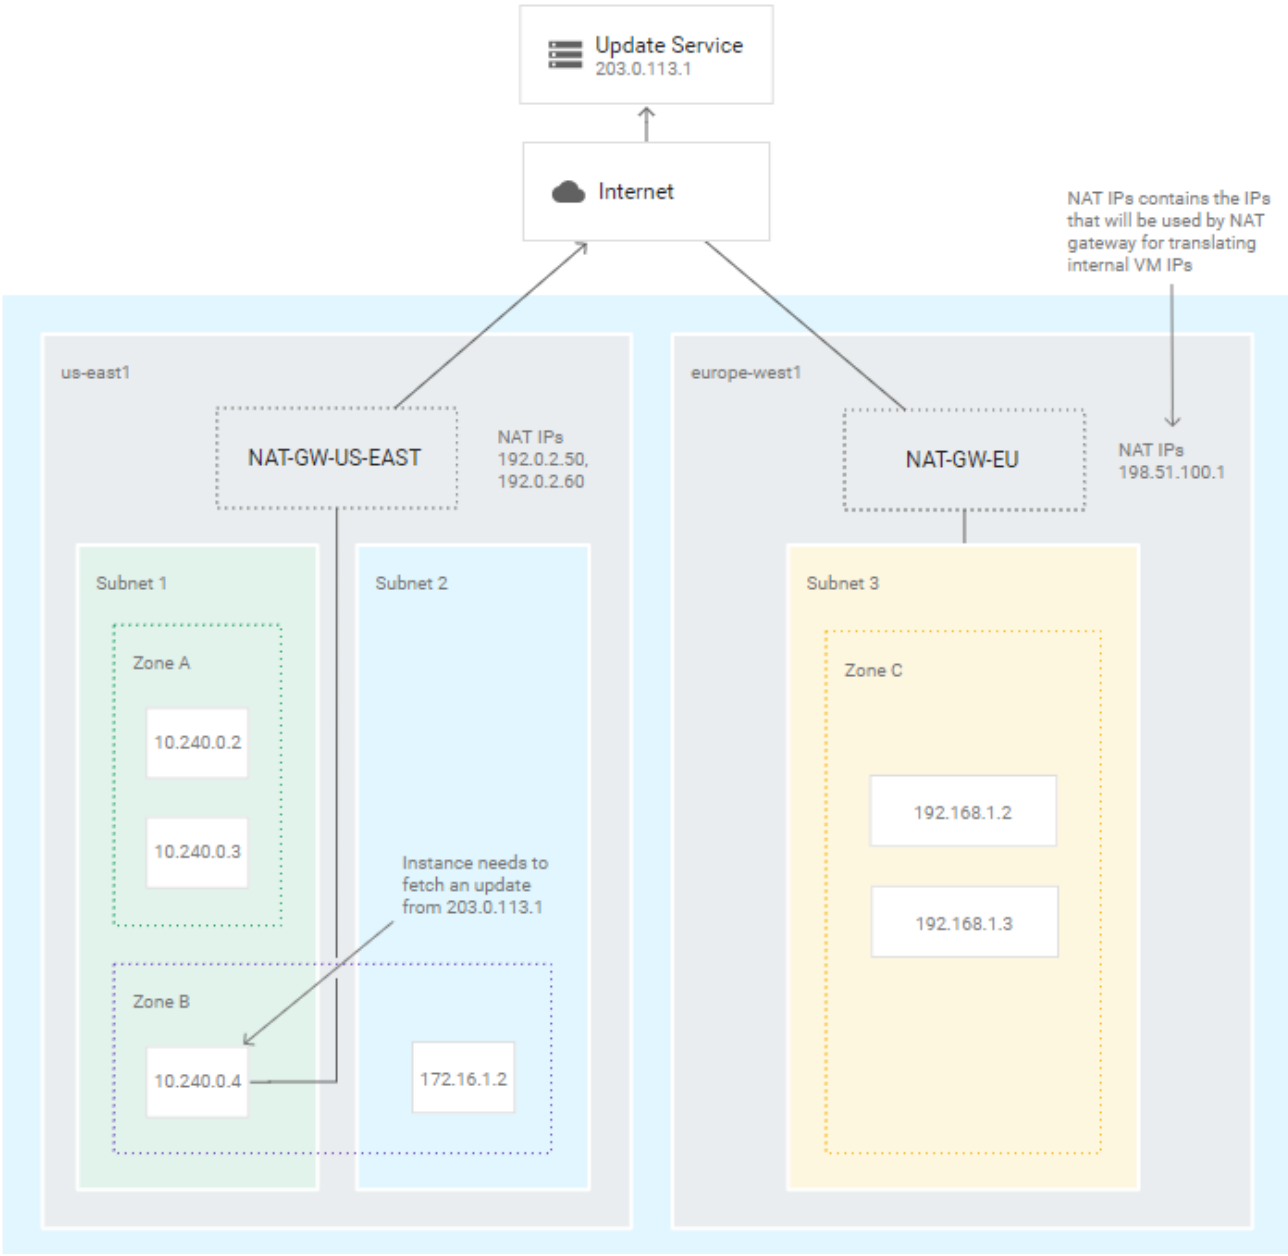
Range	Guidance
Nodes	<p>Node IPs are taken from the primary range of subnetwork associated with the cluster. Your cluster subnetwork must be large enough to fit the total number of nodes in your cluster.</p> <p>For example, if you plan to create a 900-node cluster, the subnet used with the cluster must be at least a /22 in size. A subnet size /22 contains $2^{(32-22)} = 2^{10} = 1024 - 4 \text{ reserved IP addresses} = 1020$ IP addresses, which is sufficient for the 900 node IP addresses needed for the cluster.</p>
Pods	<p>Each node currently allocates a /24 ($2^{(32-24)} = 2^8 = 256$) block of Pod IP addresses. These Pod IP addresses are taken from the associated secondary range for Pods. The Pod range as determined by the <code>--cluster-ipv4-cidr</code> or <code>--cluster-secondary-range-name</code> flags must be at least large enough to fit (total number of nodes \times 256) IP addresses.</p> <p>For example, for a 900-node cluster, you need $900 \times 256 = 230,400$ IP addresses. The IP addresses must come in /24-sized blocks, as that is the granularity assigned to a node. You need a secondary range of size /14 or larger. A /14 range of IP addresses results in $2^{(32-14)} = 2^{18} \approx 262k$ IP addresses.</p>
Services	<p>Every cluster needs to reserve a range of IP addresses for Kubernetes Service cluster IP addresses. The Service IP addresses are assigned from the associated secondary range for Services. You must ensure that the block of IP addresses is sufficient for the total number of Services that you anticipate to run in the cluster. You define the ranges defined using the <code>--services-ipv4-cidr</code> or <code>--services-secondary-range-name</code> flags.</p> <p>For example, for a cluster that runs at most 3000 Services, you need 3000 IP addresses to be used for cluster IP addresses. You need a secondary range of size /20 or larger. A /20 range of IP addresses results in $2^{(32-20)} = 2^{12} \approx 4k$ IP addresses.</p>

Kubernetes networking example

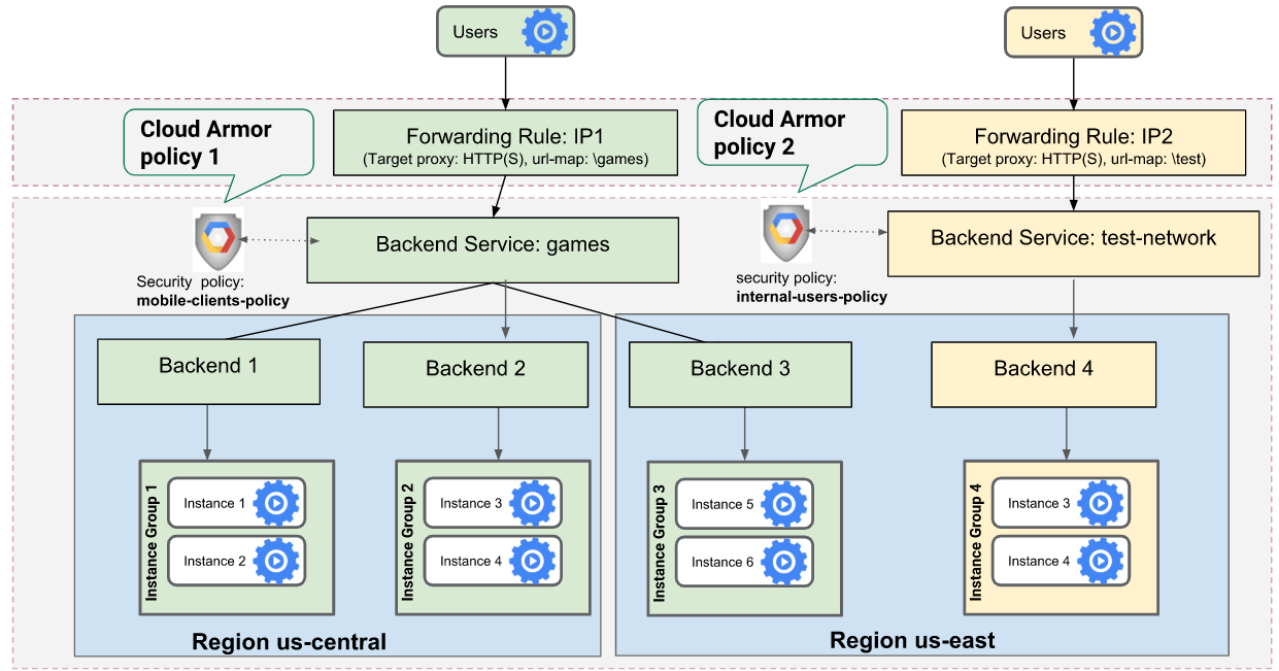


<p>Cloud DNS</p> 	<p>Session Affinity</p> 	<p>Logging</p> 	<p>Flow logs</p> 	<p>Route based VPN</p> 	<p>Policy based routing</p> 	<p>Review documents</p> <p>DNS Troubleshooting health HTTPS logging Network and Tunnel routing</p>
<p>What it is</p> <p>Cloud DNS is a high-performance, resilient, global Domain Name System (DNS) service that publishes your domain names to the global DNS in a cost-effective way.</p>	<p>What it is</p> <p>Session affinity sends all requests from the same client to the VM instance as long as the instance stays healthy and has capacity.</p>	<p>What it is</p> <p>Log generation can be exported to Stackdriver Logging, Cloud Pub/Sub, Cloud Storage, or BigQuery.</p>	<p>What it is</p> <p>VPC Flow Logs record a sample of network flows sent from and to by VM instances. These are used for monitoring, forensics, real-time security analysis, and expense optimization.</p>	<p>What it is</p> <p>Consider when the peer gateway cannot use BGP. In route based VPN, you specify only the remote traffic selector. Classic VPN static routing.</p>	<p>What it is</p> <p>Consider when the peer gateway cannot use BGP. Policy based routing uses local and a remote traffic selectors. Use with classic VPN static routing</p>	<p>Session affinity Flow Logs Import Record-sets Static Routing Labs High Throughput VPN VPC flow Logs</p>
<p>What you should know</p> <p>1- Types Zones (managed, Public, Private, forwarding, peering) 2-Internal DNS, delegated subzones)</p>	<p>What you should know</p> <p>1- Supported by the following LB (Internal, TCP and SSL proxy, HTTP(s) & Network)</p>	<p>What you should know</p> <p>1- Benefit of logging 2 – What kinds of logs 3- How to view data</p>	<p>What you should know</p> <p>1- Cases to use this to gather info to lock down access, see traffic etc 2- How to enable</p>	<p>What you should know</p> <p>1 - Local and remote traffic selector always 0.0.0.0/0</p>	<p>What you should know</p> <p>1- Configurable remote and local traffic selectors</p>	<p>My experience</p> <p>Routes based, Policy based, logging, DNS, session affinity these can be troublesome if you do not understand each concept clearly for exams. Get my hint, really troublesome.</p>
<p>Key Points</p> <p>1- On Prem connection 2- Private Zones 3- Importing Zone record-sets BIND or YAML 4- Forwarding zones</p>	<p>Key Points</p> <p>1- How each type of LB handle session affinity options (None, IP, Protocol, Port)</p>	<p>Key Points</p> <p>1- General awareness of log types and viewing.</p>	<p>Key Points</p> <p>1- What it records, how to read it</p>	<p>Key Points</p> <p>1- You must manually create and maintain the routes to the subnets in your VPC network on your peer routers.</p>	<p>Key Points</p> <p>1- You must manually create and maintain the routes to the subnets in your VPC network on your peer routers.</p>	
<p>Compute instance</p> 	<p>Key Management</p> 	<p>DNSSEC</p> 	<p>Cloud Armour</p> 	<p>Cloud NAT</p> 	<p>IKEv1/v2</p> 	<p>Review documents</p> <p>DNSSEC Rolling Updates URL_Map NAT Connecting using advanced methods</p>
<p>What it is</p> <p>Your virtual machine in the cloud. This is part of Google IaaS offering</p>	<p>What it is</p> <p>By creating and managing SSH keys, you can allow users to access a Linux instance through third-party tools.</p>	<p>What it is</p> <p>Prevents attackers from manipulating or poisoning the responses to DNS requests.</p>	<p>What it is</p> <p>Google Cloud Armor security policies are made up of rules that allow or prohibit traffic from IP addresses or ranges defined in the rule.</p>	<p>What it is</p> <p>Allows virtual machine (VM) instances without external IP addresses and private (GKE) clusters to connect to the Internet.</p>	<p>What it is</p> <p>IKEv1 limits remote traffic selectors to a single CIDR. Cloud VPN does not support creating a tunnel using IKEv1 with multiple Child SAs, each with a single CIDR</p>	<p>OS Login Video DDoS Labs HTTP Load balancer and cloud Armour</p>
<p>What you should know</p> <p>1- IP assignment internal, external 2- Static IP's</p>	<p>What you should know</p> <p>1- How to configure 2- What are the risk</p>	<p>What you should know</p> <p>1- How to set up 2- How to disable</p>	<p>What you should know</p> <p>1- Where and how it works (Edge, HTTPS load balancing proxy)</p>	<p>What you should know</p> <p>1- How it works</p>	<p>What you should know</p> <p>1- Difference between IKEv1 and IKEv2</p>	<p>My experience</p> <p>All these area combined made for some very challenging questions. It's worth spending a bit of time reviewing.</p>
<p>Key Points</p> <p>1- Testing updates 2- Tagging 3- SSH into VM</p>	<p>Key Points</p> <p>1- How to assign to your VM's 2- How to remove from VM</p>	<p>Key Points</p> <p>1- The components to make this work and be removed.</p>	<p>Key Points</p> <p>1- How it works (whitelist, blacklist) 2- DDOS</p>	<p>Key Points</p> <p>1- Hide internal IP from external host.</p>	<p>Key Points</p> <p>1- How this affect Multiple CIDR's traffic selectors Cloud VPN</p>	

NAT image




Cloud Armour image
HTTP(S) LB + Cloud Armor Data Model



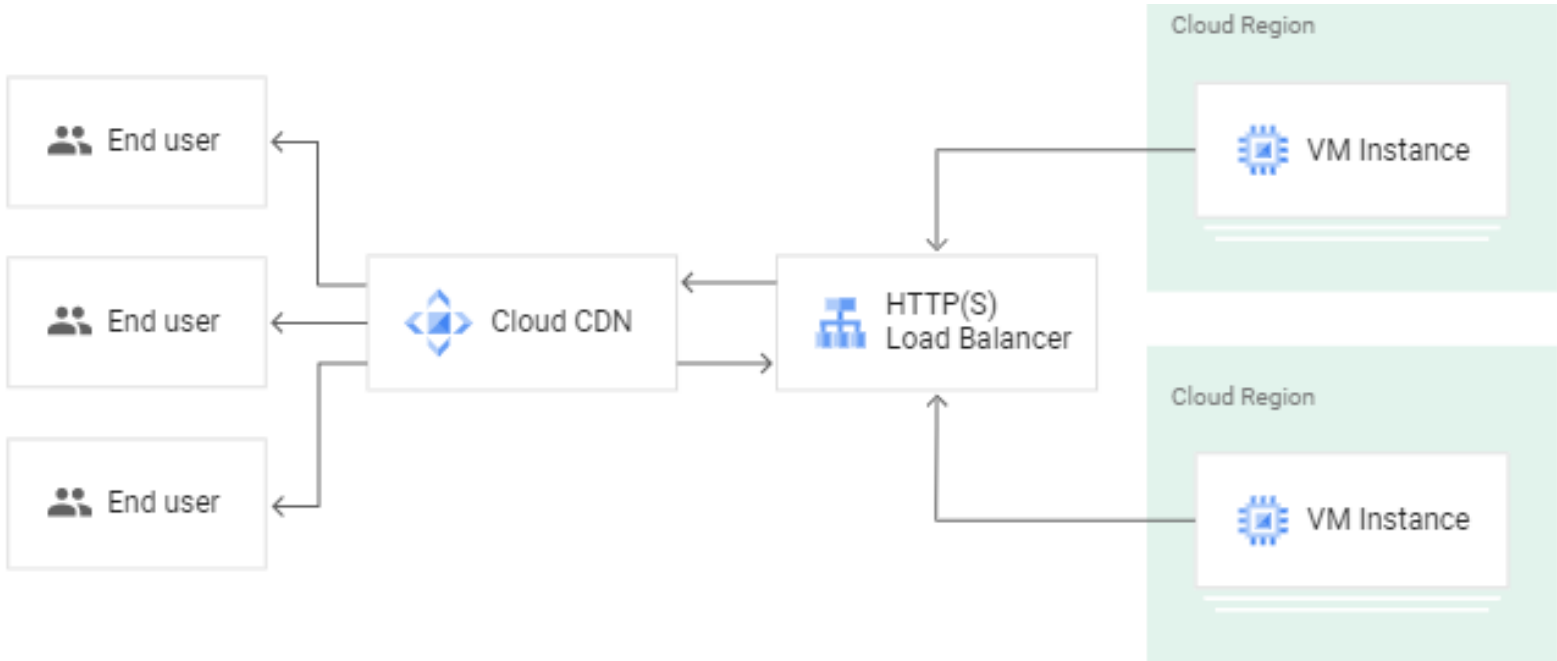
Session affinity image

Load balancer	Session affinity options
<ul style="list-style-type: none">Internal	<ul style="list-style-type: none">NoneClient IPClient IP and protocolClient IP protocol and port
<ul style="list-style-type: none">TCP ProxySSL Proxy	<ul style="list-style-type: none">NoneClient IP
<ul style="list-style-type: none">HTTP(S)	<ul style="list-style-type: none">NoneClient IPGenerated cookie
<ul style="list-style-type: none">Network	Network Load Balancing doesn't use backend services. Instead, you set session affinity for network load balancers through target pools. See the <code>sessionAffinity</code> parameter in Target Pools .

VPC Sharing 	What it is Used to connect to a common VPC network. Resources in those projects can communicate with each other securely and efficiently across project boundaries using internal IPs.	Key points 1- Centralised management 2- Firewall control 3- Internal RFC1918 communication	What you should know 1- When to use (depend of services and controls necessary etc) 2- Who gets billed	Review documents Hybrid Connectivity Shared VPC	Video CONNECTIVITY Connecting to Datacentre	My experience This will pop up. Who knows peering is sharing 😊. Core topic
VPC Peering 	What it is Access G Suite and Google Cloud features over VPN or the internet, while cutting egress fees. Connect directly with Direct Peering, or choose a partner with Carrier Peering.	Key points 1- When to peer 2- What services you have access to	What you should know 1- How to peer to a shared VPC	Review documents VPC Peering		My experience This will come. Know requirements of peering and how to peer to shared networks. Core topic
VPN 	What it is Connect your on-premises or other public cloud networks to GCP Virtual Private Cloud (VPC) securely over the internet through IPSec VPN	Key points 1- How to setup 2- IPSEC used 3- Best practices	What you should know 1- Multiple tunnels 2- ECMP	Review documents Cloud VPN		My experience Core area. Make sure you know VPN very well. Know high availability, multi tunnelling various scenarios for use.
Dedicated Interconnect 	What it is Use dedicated Interconnect to connect to Google's network through a highly available, low latency connection. (10GB higher)	Key points 1- Single mode fiber 10GBase-LR 2- LACP for links & 802.1q Vlan 3- Support EBGP with multihop 4-Ipv4 link local addresses 5- Meet at Co Location facilities	What you should know 1- Type (system and custom) 2- Default route & Subnet route 3- Static and Dynamic routes 4- Min 10GB 5- Layer2	Review documents Dedicated Interconnect		My experience Core area well represented in exam. Did I say well represented?
Partner Connect 	What it is Use Google Cloud Interconnect - Partner (Partner Interconnect) to connect to Google through a supported service provider. (from 50 MB up)	Key points 1- Best case use 2- Min size 50MB 3- Not over the internet	What you should know 1- How to setup (VLAN, Key, request location and capacity) 2- Difference L2 and L3	Review documents Partner Interconnect		My experience Core area well represented in exam also. If you don't know all the interconnect option well don't do the exam.
VLAN 	What it is VLAN attachments (also known as Interconnect Attachments) determine which Virtual Private Cloud networks can reach your on-premises network through an interconnect	Key points 1- Works with Cloud router 2- Maximum speed 10 Gbps 3- Multiple VLANs	What you should know 1- Create VLAN attachments over Cloud Interconnect connections that have passed all tests and that are ready to use	Review documents Creating VLAN attachment	Video Stackdriver	My experience Questions on this point may appears. You need a VLAN for what?
Dynamic routing 	What it is Dynamic routing is suitable for any size network. It frees you from maintaining static routes. Also, if a link fails, dynamic routing can automatically reroute traffic if possible.	Key points 1- Cloud router necessary 2- BGP session necessary	What you should know 1- IP automatically updated and propagated 2- Modes are Global or regional	Review documents Setting the network dynamic routing mode		My experience How are routes updated? Manually or automatically. Understand how this works.
Operations (stackdriver) 	What it is Stackdriver Logging allows you to store, search, analyze, monitor, and alert on log data and events from Google Cloud Platform and Amazon Web Services (AWS).	Key points 1- Individually enabled 2- Logging is supported for TCP and UDP only	What you should know 1- Troubleshooting viewing (Log entries missing, cannot view logs, where to apply logs)	Review documents Stackdriver		My experience You should have an idea where to look, what rules are logged, priorities and how to fix.

<div>Cloud CDN</div> <div></div>	<div>What it is</div> <div>Cloud CDN uses Google's global edge network to serve content closer to users, which accelerates your websites and applications.</div>	<div>Key points</div> <div>1- What it does how to enable 2- How to enable (HTTPS LB)</div>	<div>What you should know</div> <div>1- How to trouble shoot 2- Invalidation 3- Serve none cached content 4- Cache Control 5- How to enable</div>	<div>Review documents</div> <div>CDN Overview CDN Invalidation CDN troubleshooting</div>	<div>Video</div> <div>CDN Labs Cloud CDN</div>	<div>My experience</div> <div>Core area. This helps serve content faster. Know how it works well for exam.</div>
--	--	--	---	--	--	---

CDN Flow



Connection options



Interconnect comparison
Partner vs Dedicated



Partner

Customer uses service provider to meet at a Google POP

Sub-rates are available ranging from 50M to 10G per VLAN attachment

Pay for only what you need

SLAs requires at least two VLANs (default in UI) but the VLAN data rate can be different

In a customer with multiple orgs, resources are managed at the org level



Dedicated

Customer meets Google at existing POP

One to eight 10G ports available

All VLANs are over the same physical link

SLA requires at least two 10G links and associated VLANs to the VPC

All VLANs are under the same organization


Cloud
Networking

My closing thoughts

Networking is a core component of the cloud. In fact, public cloud is based on advanced SDN networking and the internet. Whether you are using code to deploy your environment or IaaS, the end result is that you want people to connect to your apps and services. If your apps are not reachable then it makes no sense. Constructing a well-defined network is important to ensure content delivery and performance is kept at its **SLO** ☺ as much as possible.

Video
[A year in GCP networking](#)

Google presence



Thanks for reviewing

Please visit the official certification outline [HERE](#)

Official practice test [HERE](#)

ps. These are my notes and tips that helped me pass the networking exam on the second attempt this is a tough exam. Every area on the document represents a topic that has a strong probability of appearing. Google may change the exam requirements at any time so always review the outline.

The knowledge is free it just cost me some time to put together. Please share with your network who may be interested in GCP Networking or need a quick refresher on networking topics.

You can also check my other Google **prep notes** for the **Security, DevOps and Engineer** exam [HERE](#)

Bonne Journée